

# Assignment 3

## Lab: Useful Tools

Prof. Ai-Chun Pang

TA / Ping-Yu Yang, Wen-Hsin Wu, Kuang-Hui Huang

---

# Goals

- **UDP socket (file transmission)**
- **Reliable data transferring (SACK)**
- **Congestion control**

---

# SACK Protocol

## GBN + SR

# SACK case 1 (working normally)




**Window (cwnd = 4):**  
The first cwnd segments

Sender



**Transmit queue:**  
Stores a list of  
**unACKed** segments

Sends data

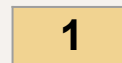


**Buffer:**  
Caches segments

Receiver



Timer

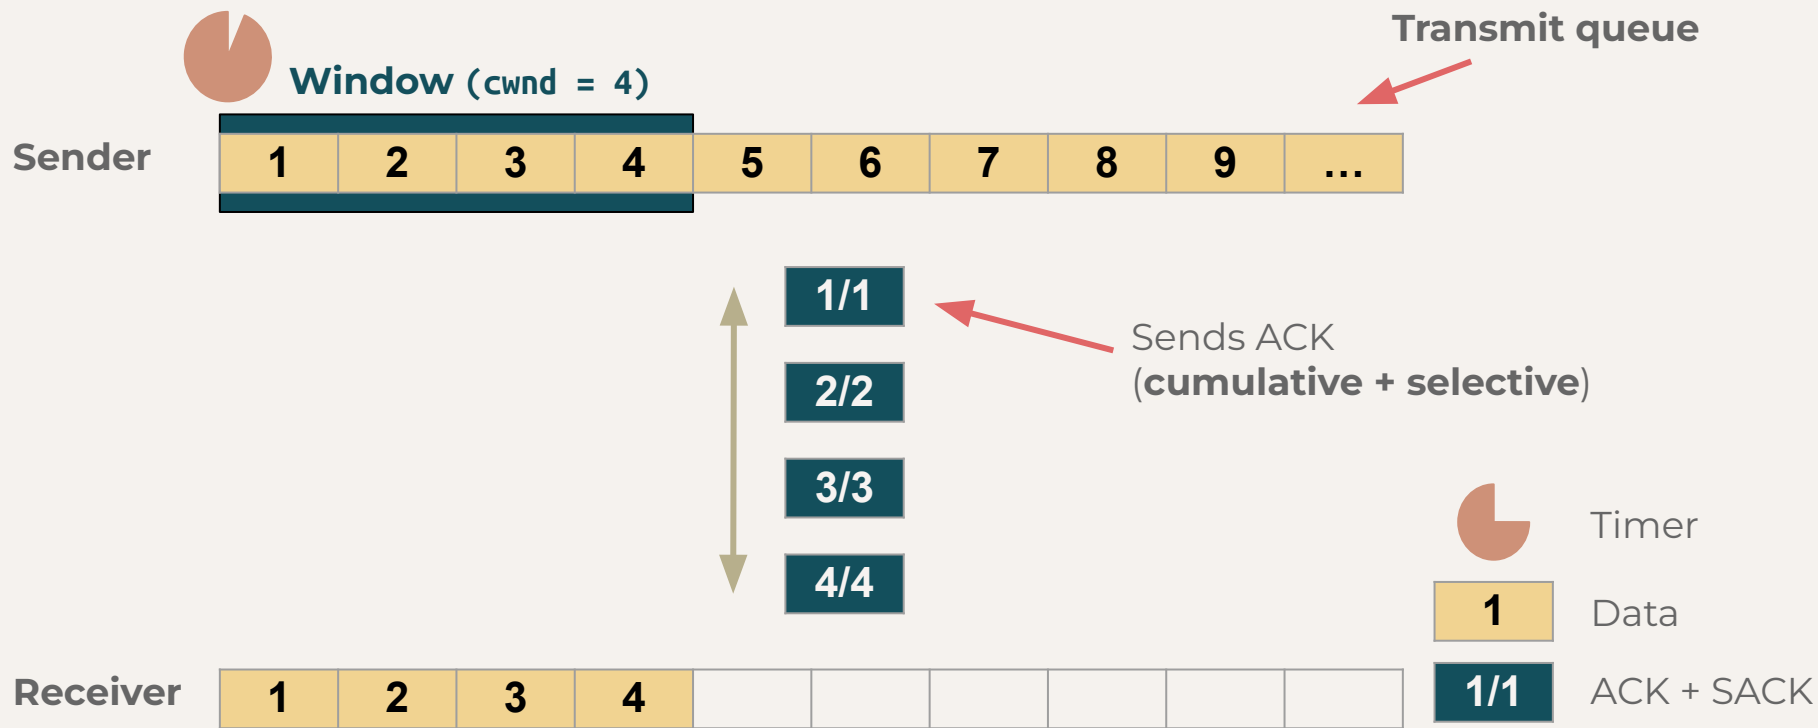


Data




ACK + SACK

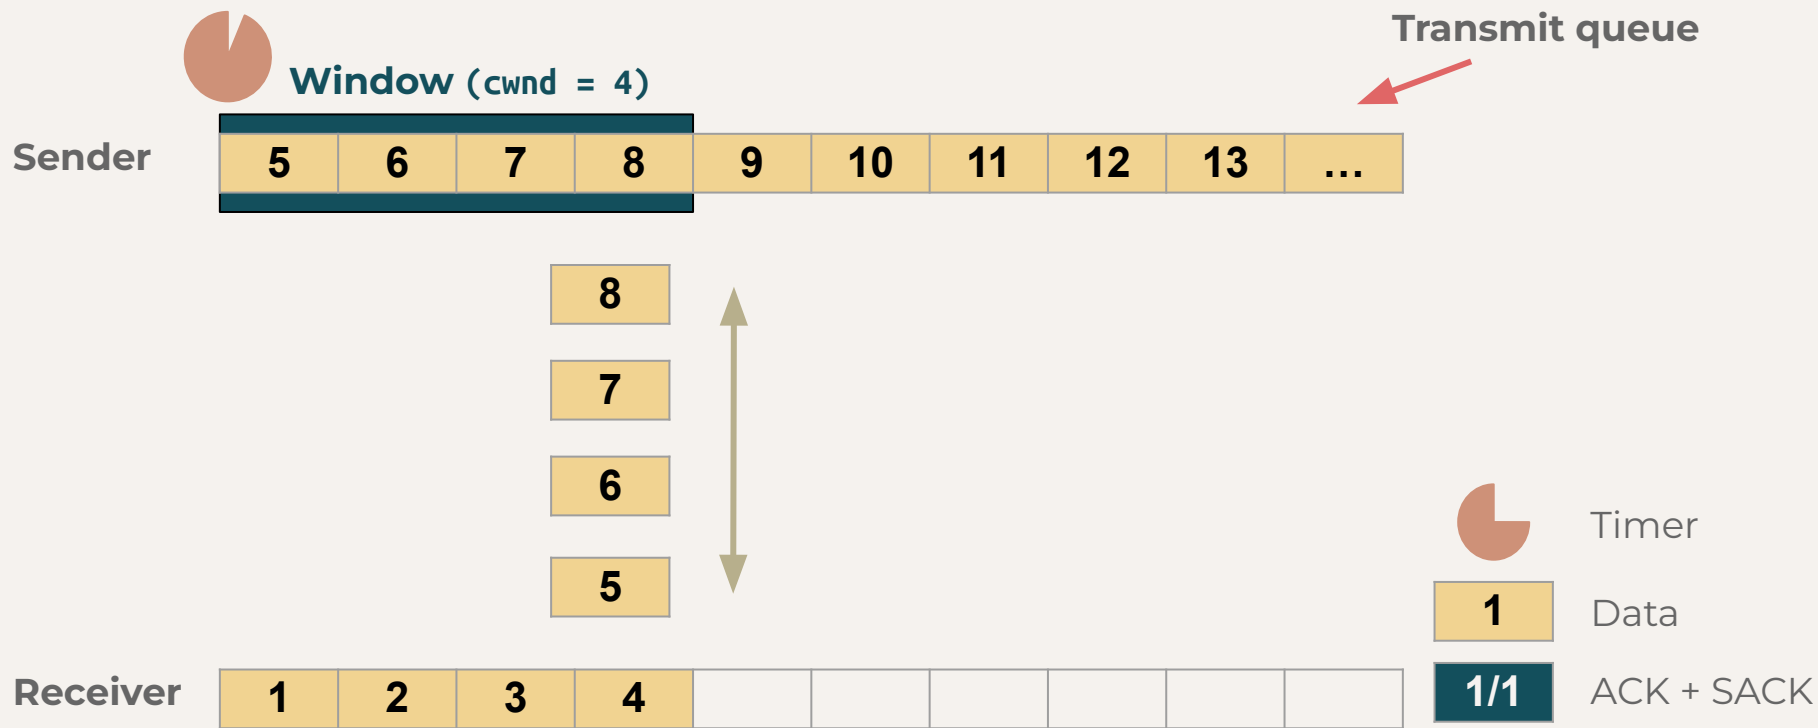
# SACK case 1 (working normally)




# SACK case 1 (working normally)



# SACK case 1 (working normally)

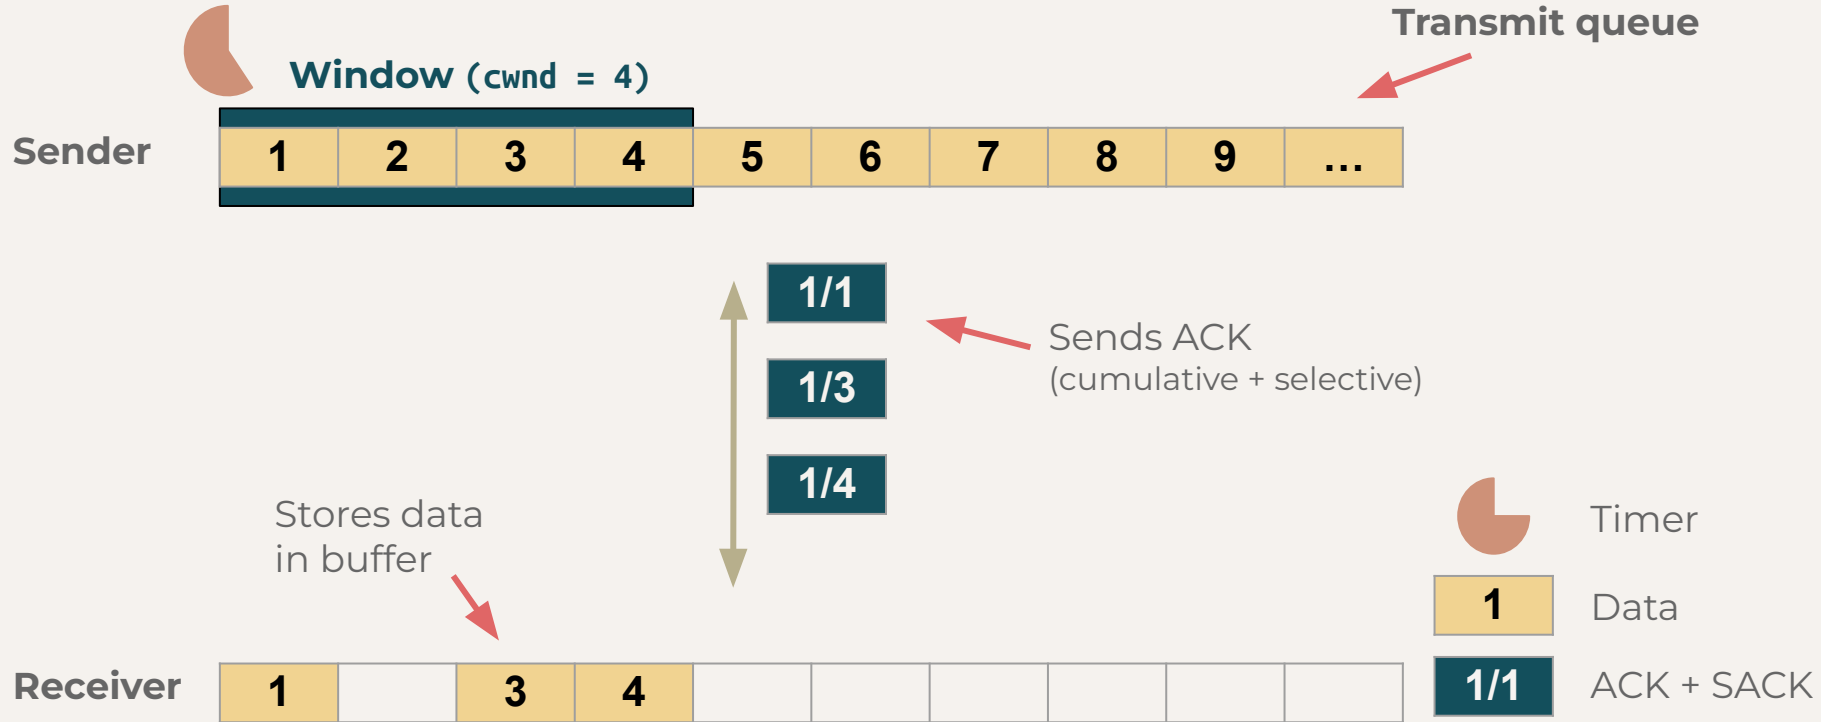


# SACK case 1 (Packet loss)

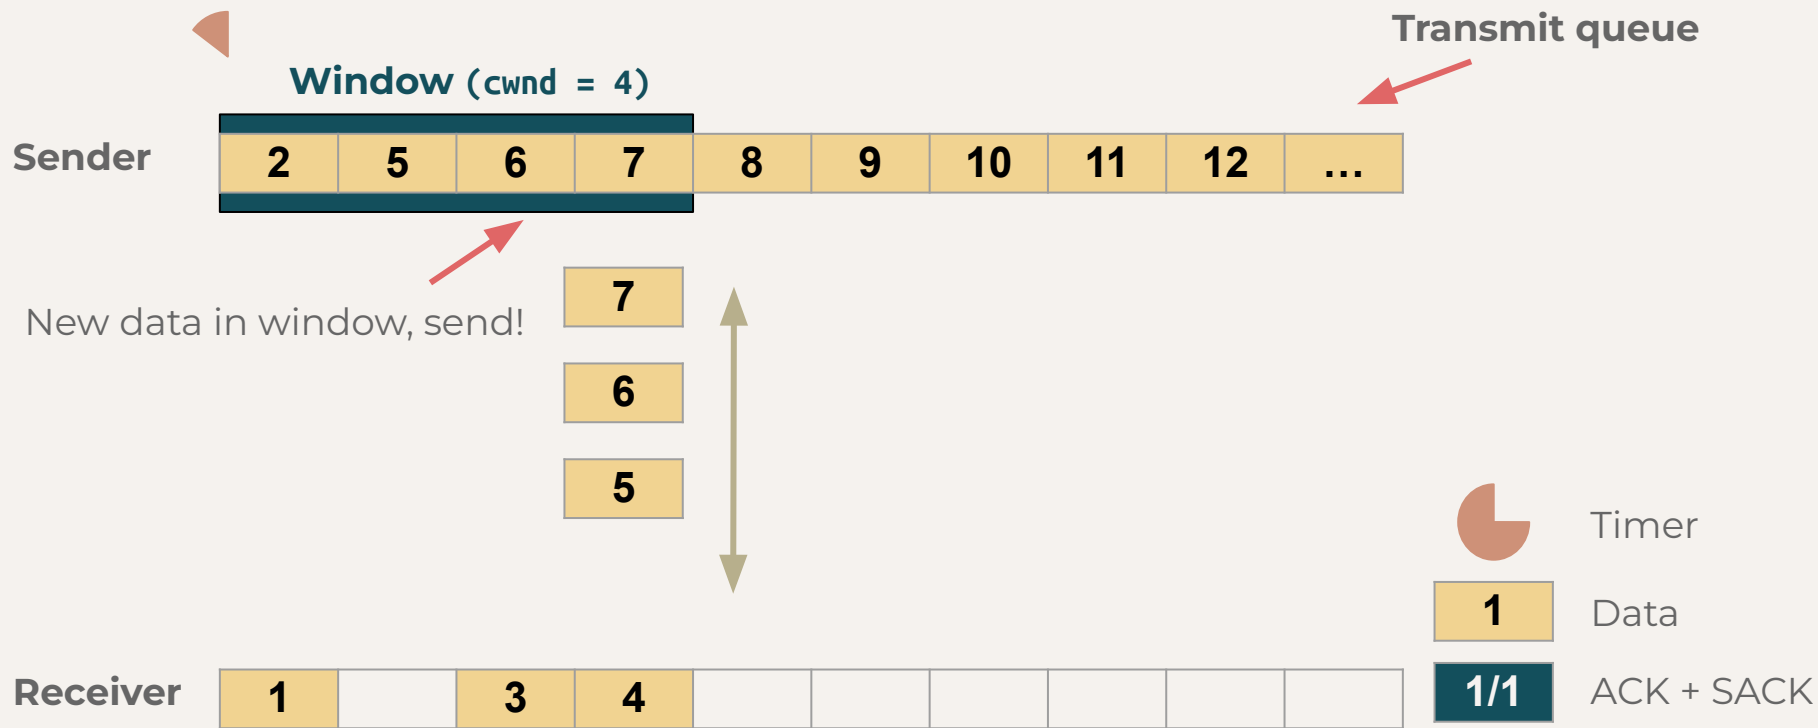




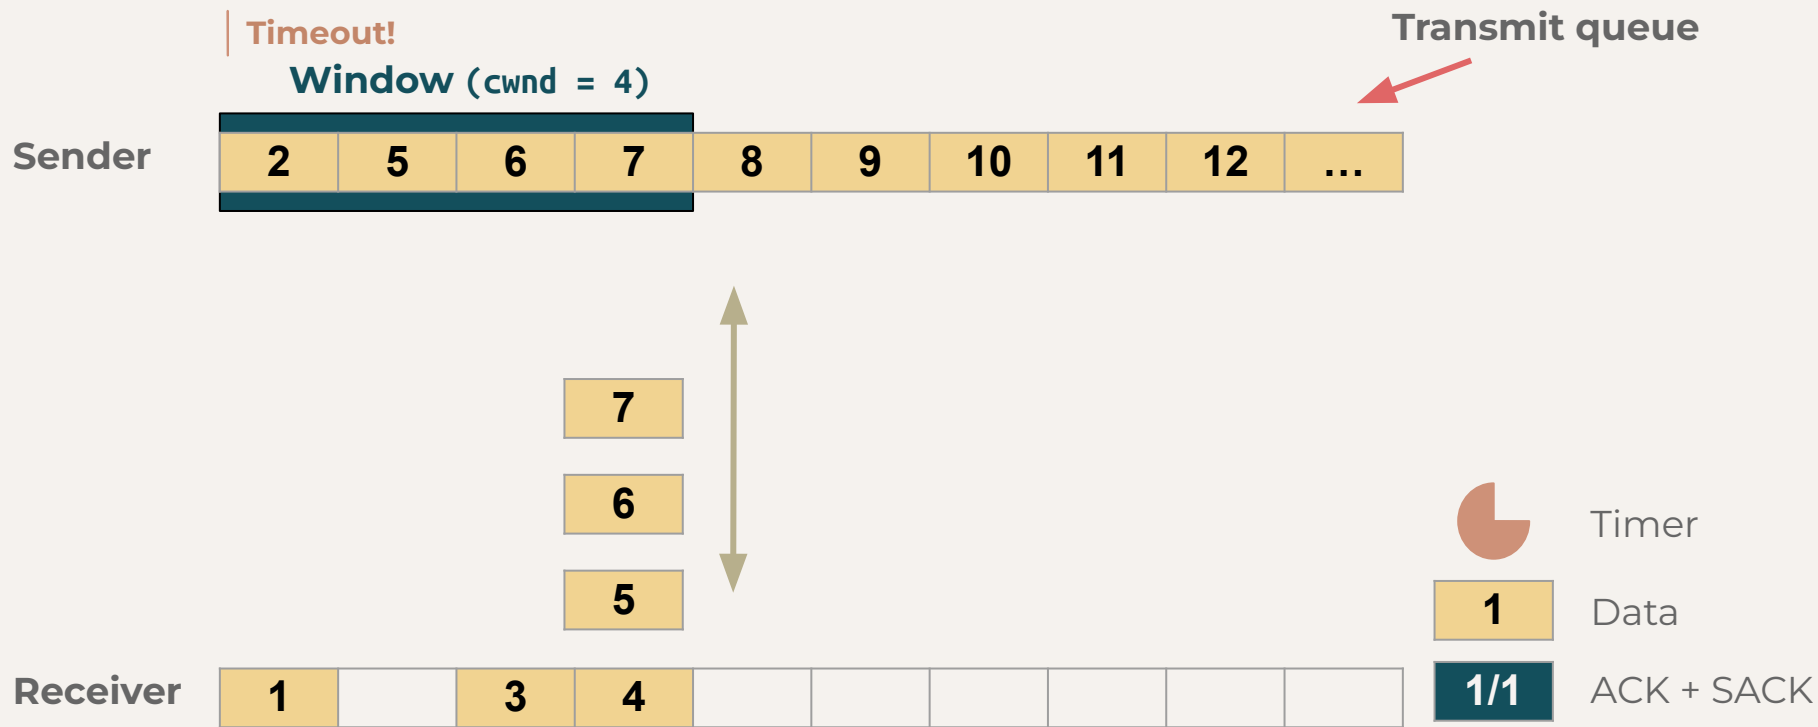
# SACK case 1 (Packet loss)



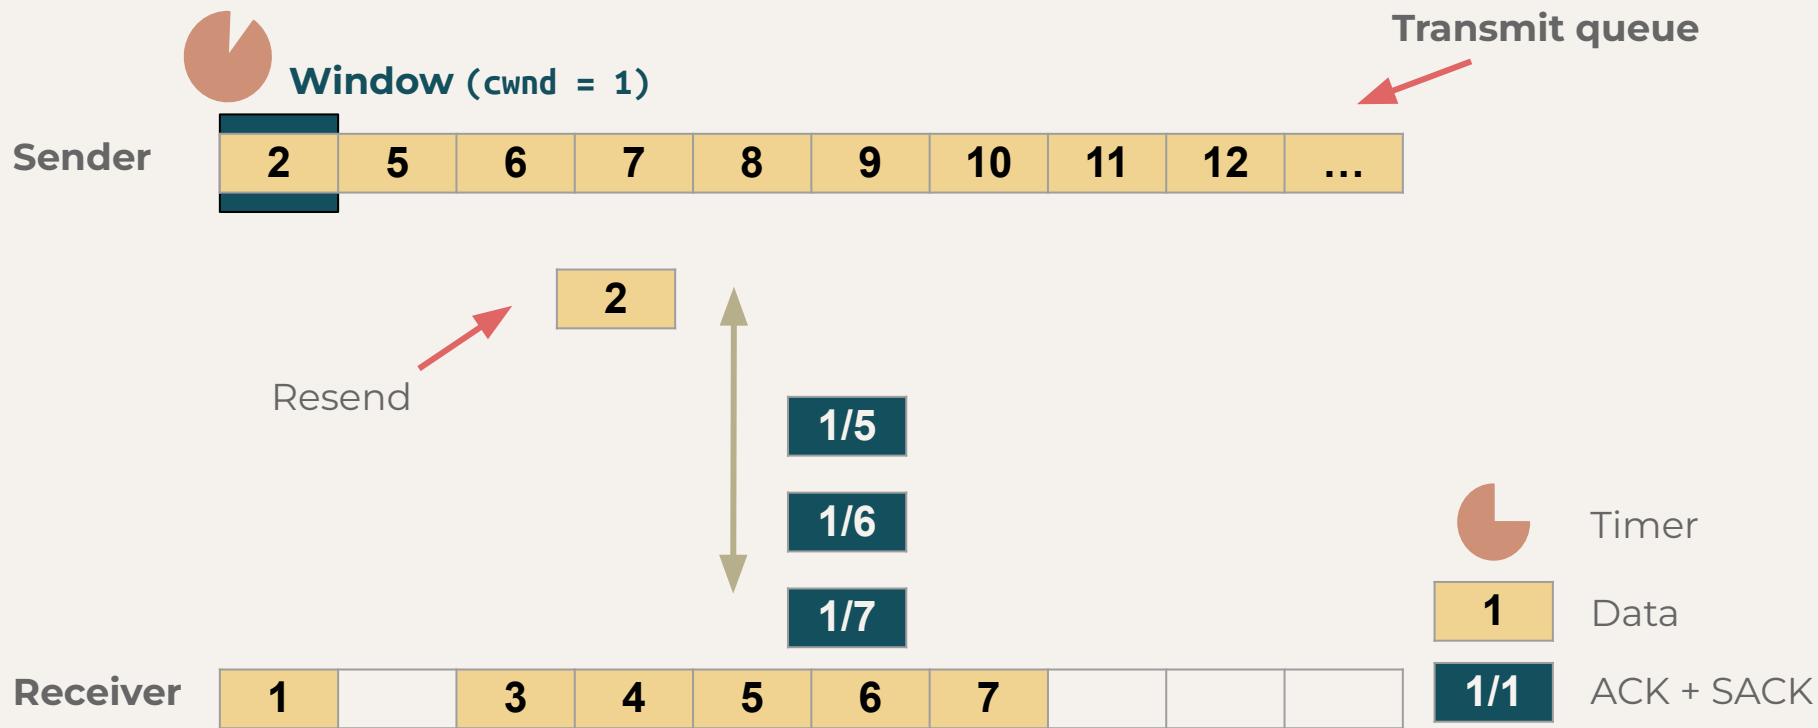
# SACK case 1 (Packet loss)



# SACK case 1 (Packet loss)



# SACK case 1 (Packet loss)




---

# SACK with Congestion Control


# SACK + Congestion Control

- Sends data 1 (cwnd = 1, threshold = 1)
- Receives ACK 1/1 (cwnd = 2, threshold = 1)




# SACK + Congestion Control

- Sends data 2, 3 (cwnd = 2, threshold = 1)
- Receives ACK 2/2, 3/3 (cwnd = 3, threshold = 1)



# SACK + Congestion Control


- Sends data 4, 5, 6 (cwnd = 3, threshold = 1)
- **Lost data 4**, receive ACK 3/5, 3/6. Duplicate cumulative ACK, cwnd = 3.






# SACK + Congestion Control

- As long as there's new segment in window, the sender will send new segments.



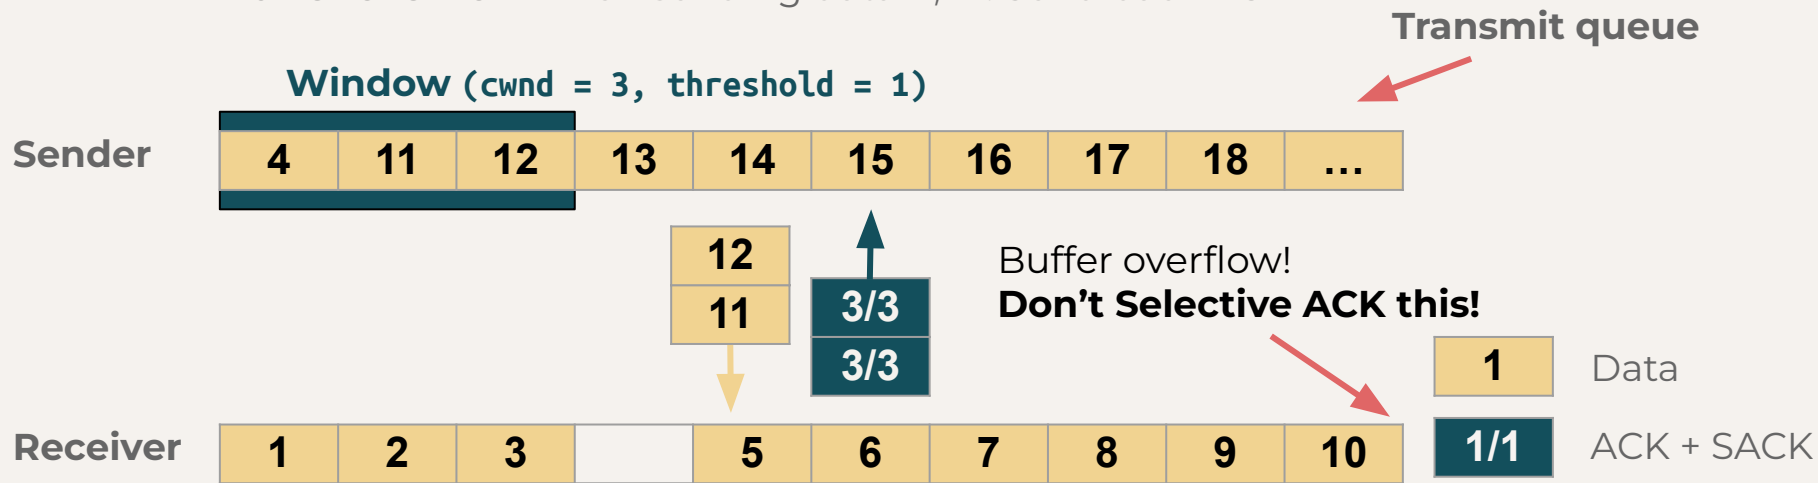
# SACK + Congestion Control

- When receiving ACK 3/7, the sender sends Data 9. But **3 duplicate cumulative ACK happened!** (ACK 3/5, 3/6, 3/7)
- To remedy packet loss, the sender **sends an additional segment (Data 4)**.




# SACK + Congestion Control

- Assuming the additional Data 4 packet is also lost, the sender will still keep sending as long as there's new segments in window.
- Buffer overflow** when sending data 11, 12. Send back ACK




# SACK + Congestion Control

- **Timeout happens!** (cwnd = 1, threshold = 1)
- Resend packet again




# SACK + Congestion Control

- After receiving Data 4, the buffer is filled.

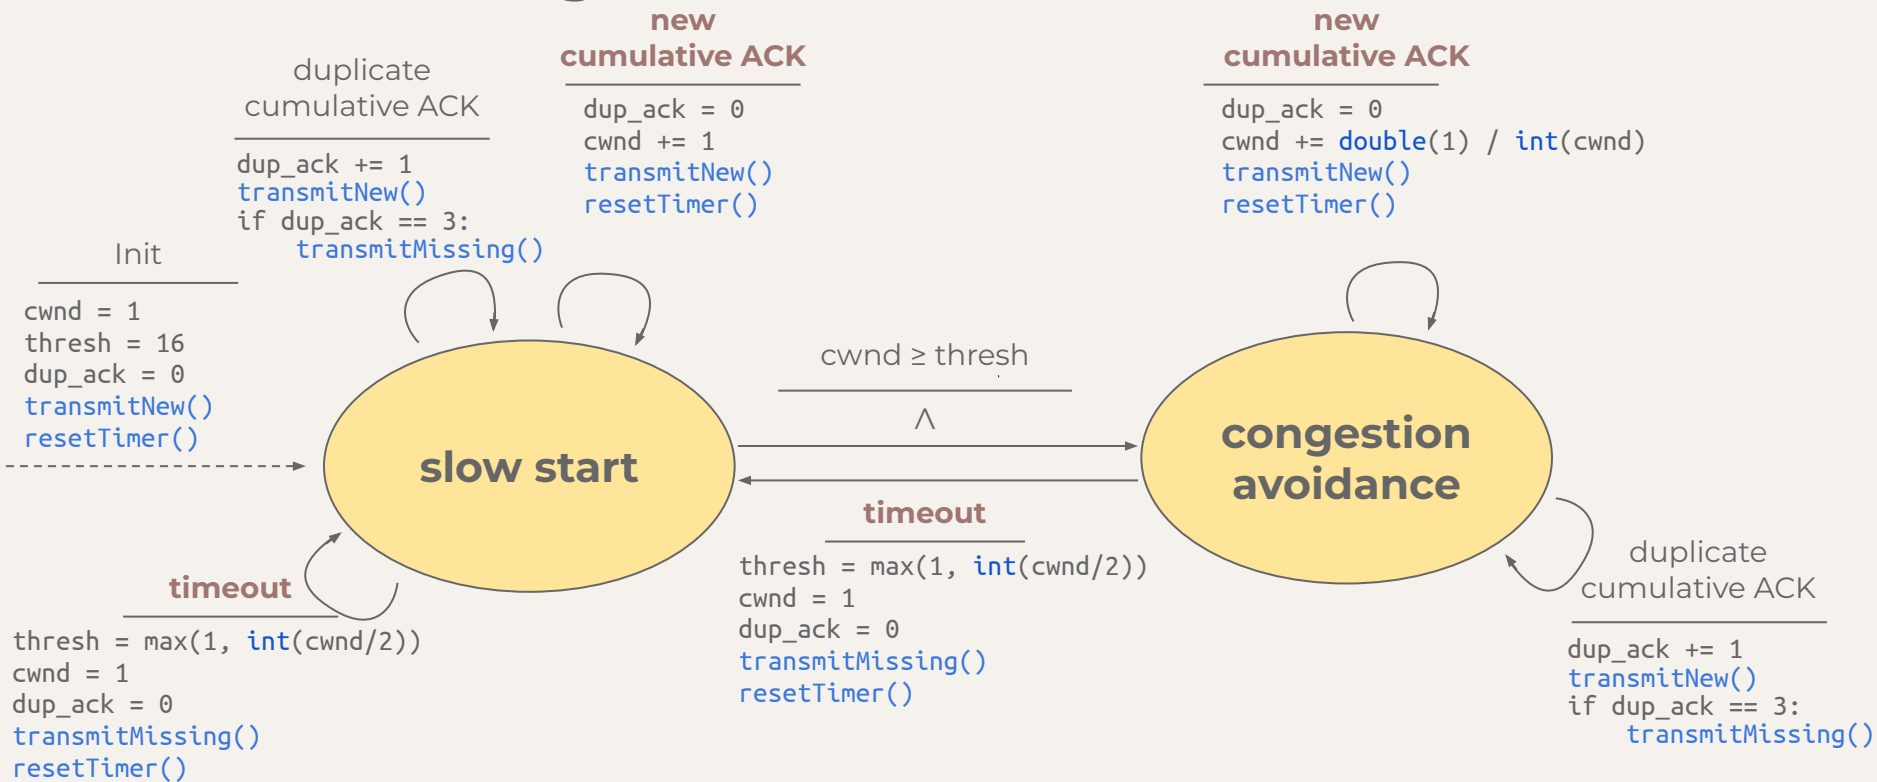


# SACK + Congestion Control

- After receiving Data 4, the buffer is filled.
- Send back ACK + flush buffer + deliver data to application



# SACK + Congestion Control



---

# Assignment 3 Announcement



# Docker

- We provide a docker config (docker-compose.yml) for you to run our example code. If you use Windows, you will need to install **Windows Subsystem Linux (WSL 2)** first.
- **Please make sure you can compile and run your code well in the provided docker container.**



# Docker Installation

- Windows
  - Install Windows Subsystem Linux (WSL): [Guide](#)
  - Install [Docker Desktop](#)
- Ubuntu
  - Install Docker through your terminal:  
# apt update  
# apt install docker.io
- macOS
  - Install [Docker Desktop](#)

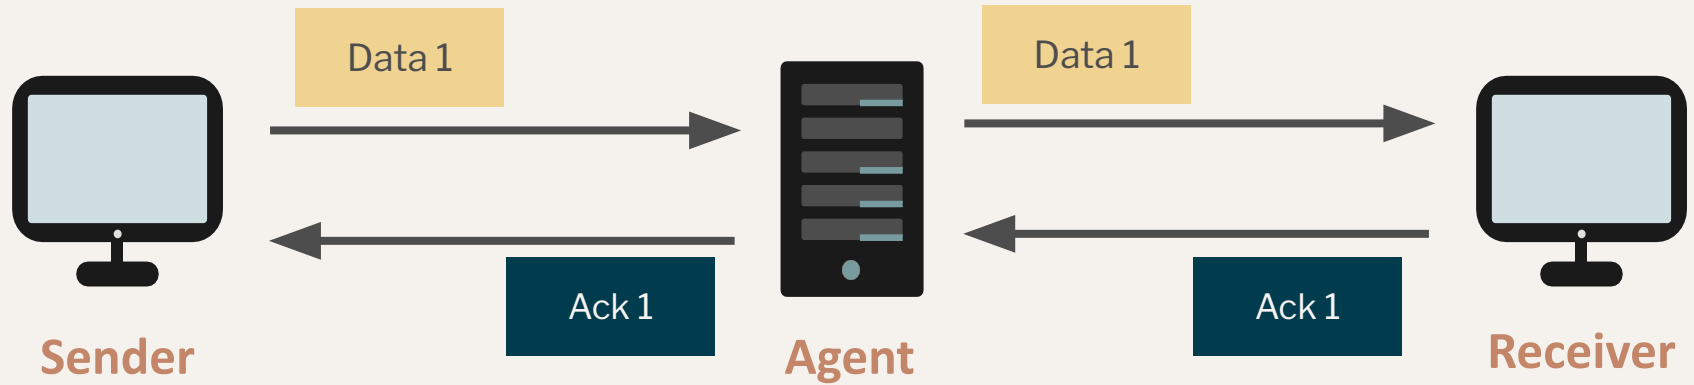
# Start Your Container

- Clone repository to your host and run the container:

```
$ git clone <your_repository>  
$ docker-compose up -d  
$ docker exec -it <container_name> bash
```

# Specification (1/16)

- Implement three components: sender, receiver and agent.



# Specification (2/16)

- **Programming language: C/C++**
- **Sender / Receiver**
  - Send / receive **file content** by UDP
  - Receiver compute **SHA-256** hash of the file stream.
  - Provide reliable transmission
  - Congestion control
- **Agent**
  - Forward Data & ACK packets
  - **Randomly drop or corrupt data packet, not ACK nor FIN**
  - Compute error rate

# Specification (3/16)

- **Reliable Transmission**
  - Data & ACK
  - Time out & Retransmission (**SACK protocol**)
  - **Sequence number**
- **Buffer handling [receiver side]**
  - Buffer Overflow: **Drop the packet** if the packet is **out of buffer**
  - Flush (write) to the file: Only when **buffer is full** or **all segments for the file are received**.

# Specification (4/16)

- **Congestion Control** [**sender side**]
  - Slow Start
    1. Send single packet in the beginning
    2. When window size is under the threshold, it increases **exponentially** until packet loses
    3. When window size is equal to or over the threshold, it increases **linearly** until packet loses
  - Packet loss / Time out
    4. Set **threshold** to  $\max(1, \text{int}(\text{cwnd} / 2))$
    5. Set **window size** to 1
    6. Retransmit – for the first “unACKed packet” (`transmitMissing()`)

# Specification (5/16)

- **Congestion Control** [**sender side**]
  - Duplicate cumulative ACK
    1. Detect when **3** duplicate cumulative ACK has happened.
    2. Retransmit the first “unACKed packet” (`transmitMissing()`)



# Specification (6/16)

- **Show Message:** Log to stdout
  - Sender:
    - send, recv, data, ack, fin, finack, sequence number, ack number, sack number, time out, resnd, winSize, threshold
  - Receiver:
    - send, recv, data, ack, fin, finack, sequence number, ack number, sack number, drop (corrupted/buffer overflow), flush, sha256, finsha
  - Agent:
    - get, fwd, data, ack, fin, finack, sequence number, ack number, sack number, drop, corrupt, error rate

# Specification (7/16)

- Show Message for **sender**

Uses tab: send\tdata\t#%,\twinSize = %d

ACK segment have  
cumulative & selective  
ACK number

```
send data #1, winSize = 1
recv ack #1, sack #1
send data #2, winSize = 2
send data #3, winSize = 2
recv ack #2, sack #2
send data #4, winSize = 3
send data #5, winSize = 3
recv ack #3, sack #3
send data #6, winSize = 4
send data #7, winSize = 4
recv ack #3, sack #5
send data #8, winSize = 4
...
send data #264, winSize = 17
recv ack #191, sack #191
recv ack #191, sack #191
time out, threshold = 8, winSize = 1
resnd data #192, winSize = 1
```

Out of order SACK-ed #5  
Data#8 gets into window, send!

Might be “send”  
or “resnd”

Timeout!  
TransmitMissing()

# Specification (8/16)

- Show Message for **sender**

Every data segments  
has been ACKed!



Finish the stream  
(agent will never drop those 2)  
(there may be more ACK in between)



```
...
resnd data #1015, winSize = 8
recv ack #1014, sack #1019
recv ack #1014, sack #1020
recv ack #1014, sack #1021
recv ack #1014, sack #1023
recv ack #1014, sack #1014
recv ack #1021, sack #1015
time out, threshold = 4, winSize = 1
resnd data #1022, winSize = 1
recv ack #1023, sack #1022
resnd data #1024, winSize = 2
recv ack #1024, sack #1024
send fin
recv finack
```



New ACK!  
But no new segments gets  
into the window...

# Specification (9/16)

- Show Message for **receiver**
- Check a packet in the following order:
  - corrupt** or not
  - in / under / above buffer**

```
[ recv  data  #1    (in order)
  send  ack   #1,   sack  #1
  recv  data  #2    (in order)
  send  ack   #2,   sack  #2
  recv  data  #3    (in order)
  send  ack   #3,   sack  #3
[ recv  data  #5    (out of order, sack-ed)
  send  ack   #3,   sack  #5
  recv  data  #6    (out of order, sack-ed)
  send  ack   #3,   sack  #6
  recv  data  #7    (out of order, sack-ed)
  send  ack   #3,   sack  #7
[ drop  data  #8    (corrupted)
  send  ack   #3,   sack  #3
  recv  data  #9    (out of order, sack-ed)
  send  ack   #3,   sack  #9
  recv  data  #10   (out of order, sack-ed)
  send  ack   #3,   sack  #10
[ recv  data  #4    (in order)
  send  ack   #7,   sack  #4
...

```

# Specification (10/16)

Corrupt but still sends SACK  
(but sackNumber = ackNumber)

```
recv data #1 (in order)
send ack #1, sack #1
recv data #2 (in order)
send ack #2, sack #2
recv data #3 (in order)
send ack #3, sack #3
recv data #5 (out of order, sack-ed)
send ack #3, sack #5
recv data #6 (out of order, sack-ed)
send ack #3, sack #6
recv data #7 (out of order, sack-ed)
send ack #3, sack #7
drop data #8 (corrupted)
send ack #3, sack #3
recv data #9 (out of order, sack-ed)
send ack #3, sack #9
recv data #10 (out of order, sack-ed)
send ack #3, sack #10
recv data #4 (in order)
send ack #7, sack #4
...
```

In order! Update base!

Out of order but in buffer range.  
Store in buffer and SACK.

In order! Update base!  
(we haven't SACKed Data#8)

# Specification (11/16)

Buffer overflow, drop & send SACK →

Received data and buffer is filled!  
Send ACK + Flush + output sha256 →

Buffer is filled, so flush + sha256 →

Received FIN, so flush + sha256 →

File stream finished, so finsha →

```
...
drop  data  #258   (buffer overflow)
send  ack    #231, sack  #231
recv  data  #232   (in order)
send  ack    #256, sack  #232
flush
sha256 256000 63ca53460cf467cfeccd33b2781b7927554a6c04d4a04eee0158286365b1d204
recv  data  #260   (out of order, sack-ed)
send  ack    #256, sack  #260
...
recv  data  #1024  (in order)
send  ack    #1024, sack  #1024
flush
sha256 1024000 20160a71cc0f658c4e78070e2629241c2e02b339bb8ca74c99e88b1fb4156ba2
recv  fin
send  finack
flush
sha256 1024000 20160a71cc0f658c4e78070e2629241c2e02b339bb8ca74c99e88b1fb4156ba2
finsha 20160a71cc0f658c4e78070e2629241c2e02b339bb8ca74c99e88b1fb4156ba2
```

# Specification (12/16)

- Show Message for **agent**

```
[ get    data    #1
  fwd    data    #1,    error rate = 0.0000
  get    ack     #1,    sack    #1
  fwd    ack     #1,    sack    #1
  get    data    #2
  fwd    data    #2,    error rate = 0.0000
  get    data    #3
  fwd    data    #3,    error rate = 0.0000
  get    ack     #2,    sack    #2
  fwd    ack     #2,    sack    #2
  get    ack     #3,    sack    #3
  fwd    ack     #3,    sack    #3
[ get    data    #4
  drop   data    #4,    error rate = 0.2500
  get    data    #5
  fwd    data    #5,    error rate = 0.2000
```

```
...
[ get    data    #1024
  corruptdata    #1024, error rate = 0.1001
  get    ack     #1014, sack    #1020
  fwd    ack     #1014, sack    #1020
  ...
  get    data    #1024
  fwd    data    #1024, error rate = 0.0998
  get    ack     #1024, sack    #1024
  fwd    ack     #1024, sack    #1024
[ get    fin
  fwd    fin
[ get    finack
  fwd    finack
```

There may be more  
get & fwd in between  
FIN & FINACK

# Specification (13/16)

- **Packet structure**

- The format used for transmission should be the same as the right side (defined in def.h):
- fin: 0 or 1
- syn: 0 or 1 (just make it 0)
- ack: 0 or 1
- checksum: we will use `crc32()` in `zlib.h` to calculate checksum

```
struct header {  
    int length;  
    int seqNumber;  
    int ackNumber;  
    int sackNumber;  
    int fin;  
    int syn;  
    int ack;  
    unsigned long checksum;  
};  
  
struct segment {  
    struct header head;  
    char data[MAX_SEG_SIZE];  
};
```



# Specification (14/16)

- **Settings**

- Sender
  - Default threshold: 16
  - Default window size: 1
- Receiver
  - Default packet data size (payload): MAX\_SEG\_SIZE (1000) bytes
  - Default buffer size: MAX\_SEG\_BUF\_SIZE (256) (# of packets)
- Agent
  - Default packet data size (payload): MAX\_SEG\_SIZE (1000) bytes
- Default time out: TIMEOUT\_MILLISECONDS (1000) milliseconds.

# Specification (15/16)

- You are required to write a Makefile for compilation.

```
$ make sender      // To compile sender code
$ make agent       // To compile agent code
$ make receiver    // To compile receiver code
$ make             // To run the above 3 compilation
```

- After compilation, there will be 3 binary files named “sender,” “agent,” and “receiver.”

# Specification (16/16)

- Execute the following commands in different terminals and **in sequence**.

```
$ ./agent <agent_port> <send_ip> <send_port> <recv_ip> <recv_port> <error_rate>
$ ./receiver <recv_ip> <recv_port> <agent_ip> <agent_port> <dst_filepath>
$ ./sender <send_ip> <send_port> <agent_ip> <agent_port> <src_filepath>
```

- The error rate will be a floating point number between 0 and 1.
- Only legal IP and port will be used in this assignment.
- `src_filepath` may be `/dev/stdin`, and `dst_filepath` may be `/dev/null`.
- The input file is a binary file, i.e., it can have null bytes.

# Grading Policy (1/2)

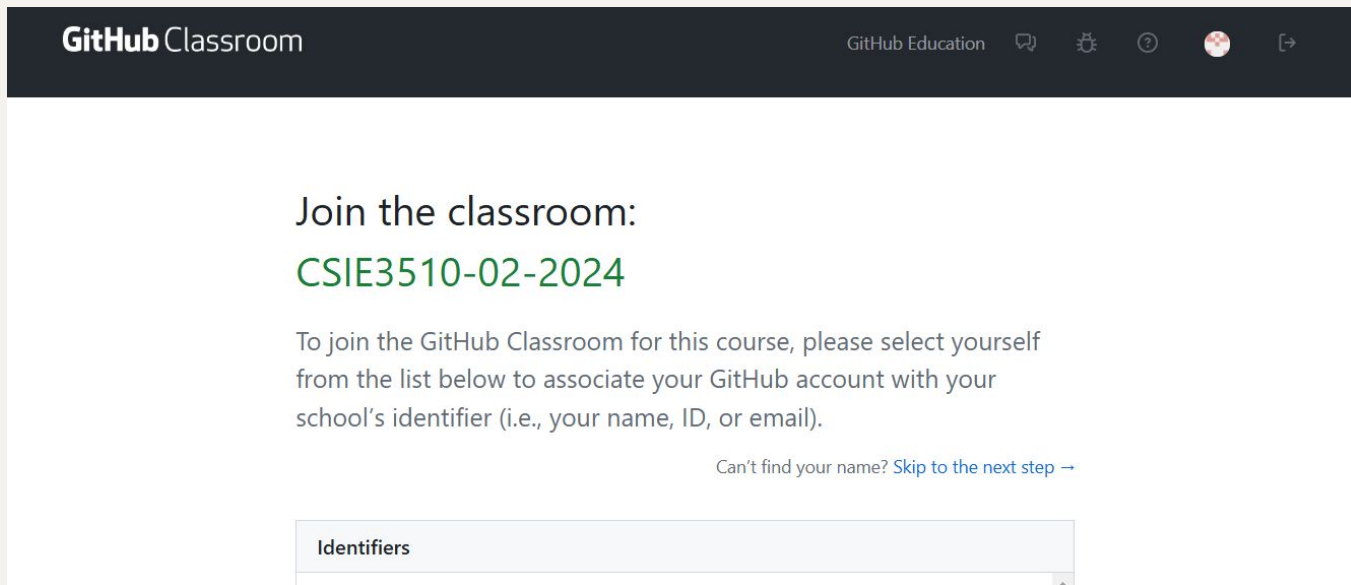
- This assignment accounts for 15% of the total score.
- **File Transmission (20%)**
  - Receives and stores the file correctly
- **Buffer handling (10%)**
  - Flush + drop buffer overflow
- **Reliable transmission (20%)**
  - SACK protocol
- **SHA-256 Hash (10%)**
  - finsha & sha256

# Grading Policy (2/2)

- **Congestion control (15%)**
  - Window size & threshold
- **Show Message (10%)**
  - Show message correctly
- **Report (15%)**
  - Explain your program structure (5% \* 3)  
(including **3 flow charts** for **sender**, **agent**, and **receiver**)

# Github Classroom

- Get the access of assignment materials via **Github Classroom**.



The screenshot shows the GitHub Classroom interface. At the top is a dark navigation bar with the 'GitHub Classroom' logo on the left and links for 'GitHub Education', a chat icon, a GitHub logo, a help icon, a profile icon, and an external link icon. The main content area is white and contains the following text:

Join the classroom:

**CSIE3510-02-2024**

To join the GitHub Classroom for this course, please select yourself from the list below to associate your GitHub account with your school's identifier (i.e., your name, ID, or email).

Can't find your name? [Skip to the next step →](#)

Below the text is a light blue box with the title 'Identifiers'. The box is currently empty, indicating a list of identifiers that has not yet been populated or displayed.

# Submission

- Report
  - Your report should be a **pdf** file. Submit it to **Gradescope**.
  - PDF file name: <studentID>\_hw3.pdf
    - e.g., B11902999\_hw3.pdf
- Codes
  - Please push all the **source code** (i.e., without your report, and the execution file) to **Github classroom** assignment.
- The penalty for the wrong format is **10 points**.
- **No plagiarism is allowed. A plagiarist will be graded zero.**

---

# Submission

- Deadline
  - Due Date : 13:00, December 11<sup>st</sup>, 2024 (temporary)
  - Penalty for late submission is **20 points per day**.



# Sample Codes

- We will provide sample codes for your reference
  - *sender.cpp, receiver.cpp, agent.cpp*
  - *crc32.cpp, sha256.cpp*
  - *def.h*
  - *Makefile*

---

# Supplementary Materials

## crc32()

- Compute CRC-32 Checksum. (ref. `crc32.cpp`)

```
#include <zlib.h>
unsigned long crc32(unsigned long crc, const Bytef * buf, unsigned int len);
```

- **crc**
  - The previous value for the checksum.
  - In this homework, we can set it to 0L.
- **buf**: Specifies the buffer to contain the data to be added to this checksum.
- **len**: Specifies the size of `buf`.

# SHA-256

- Compute SHA-256 hash. (ref. sha256.cpp)

```
#include <openssl/evp.h>
EVP_MD_CTX *EVP_MD_CTX_new(void);
int EVP_DigestInit_ex(EVP_MD_CTX *ctx, const EVP_MD *type, ENGINE *impl);
int EVP_DigestUpdate(EVP_MD_CTX *ctx, const void *d, size_t cnt);
int EVP_MD_CTX_copy_ex(EVP_MD_CTX *out, const EVP_MD_CTX *in);
int EVP_DigestFinal_ex(EVP_MD_CTX *ctx, unsigned char *md, unsigned int *s);
void EVP_MD_CTX_free(EVP_MD_CTX *ctx);
```

---

Contact us if you have any problem. ●ω●)๓

TA Email: [ntu.cnta@gmail.com](mailto:ntu.cnta@gmail.com)